FlagQuantum distributed differentiable statevector · same-host scaling

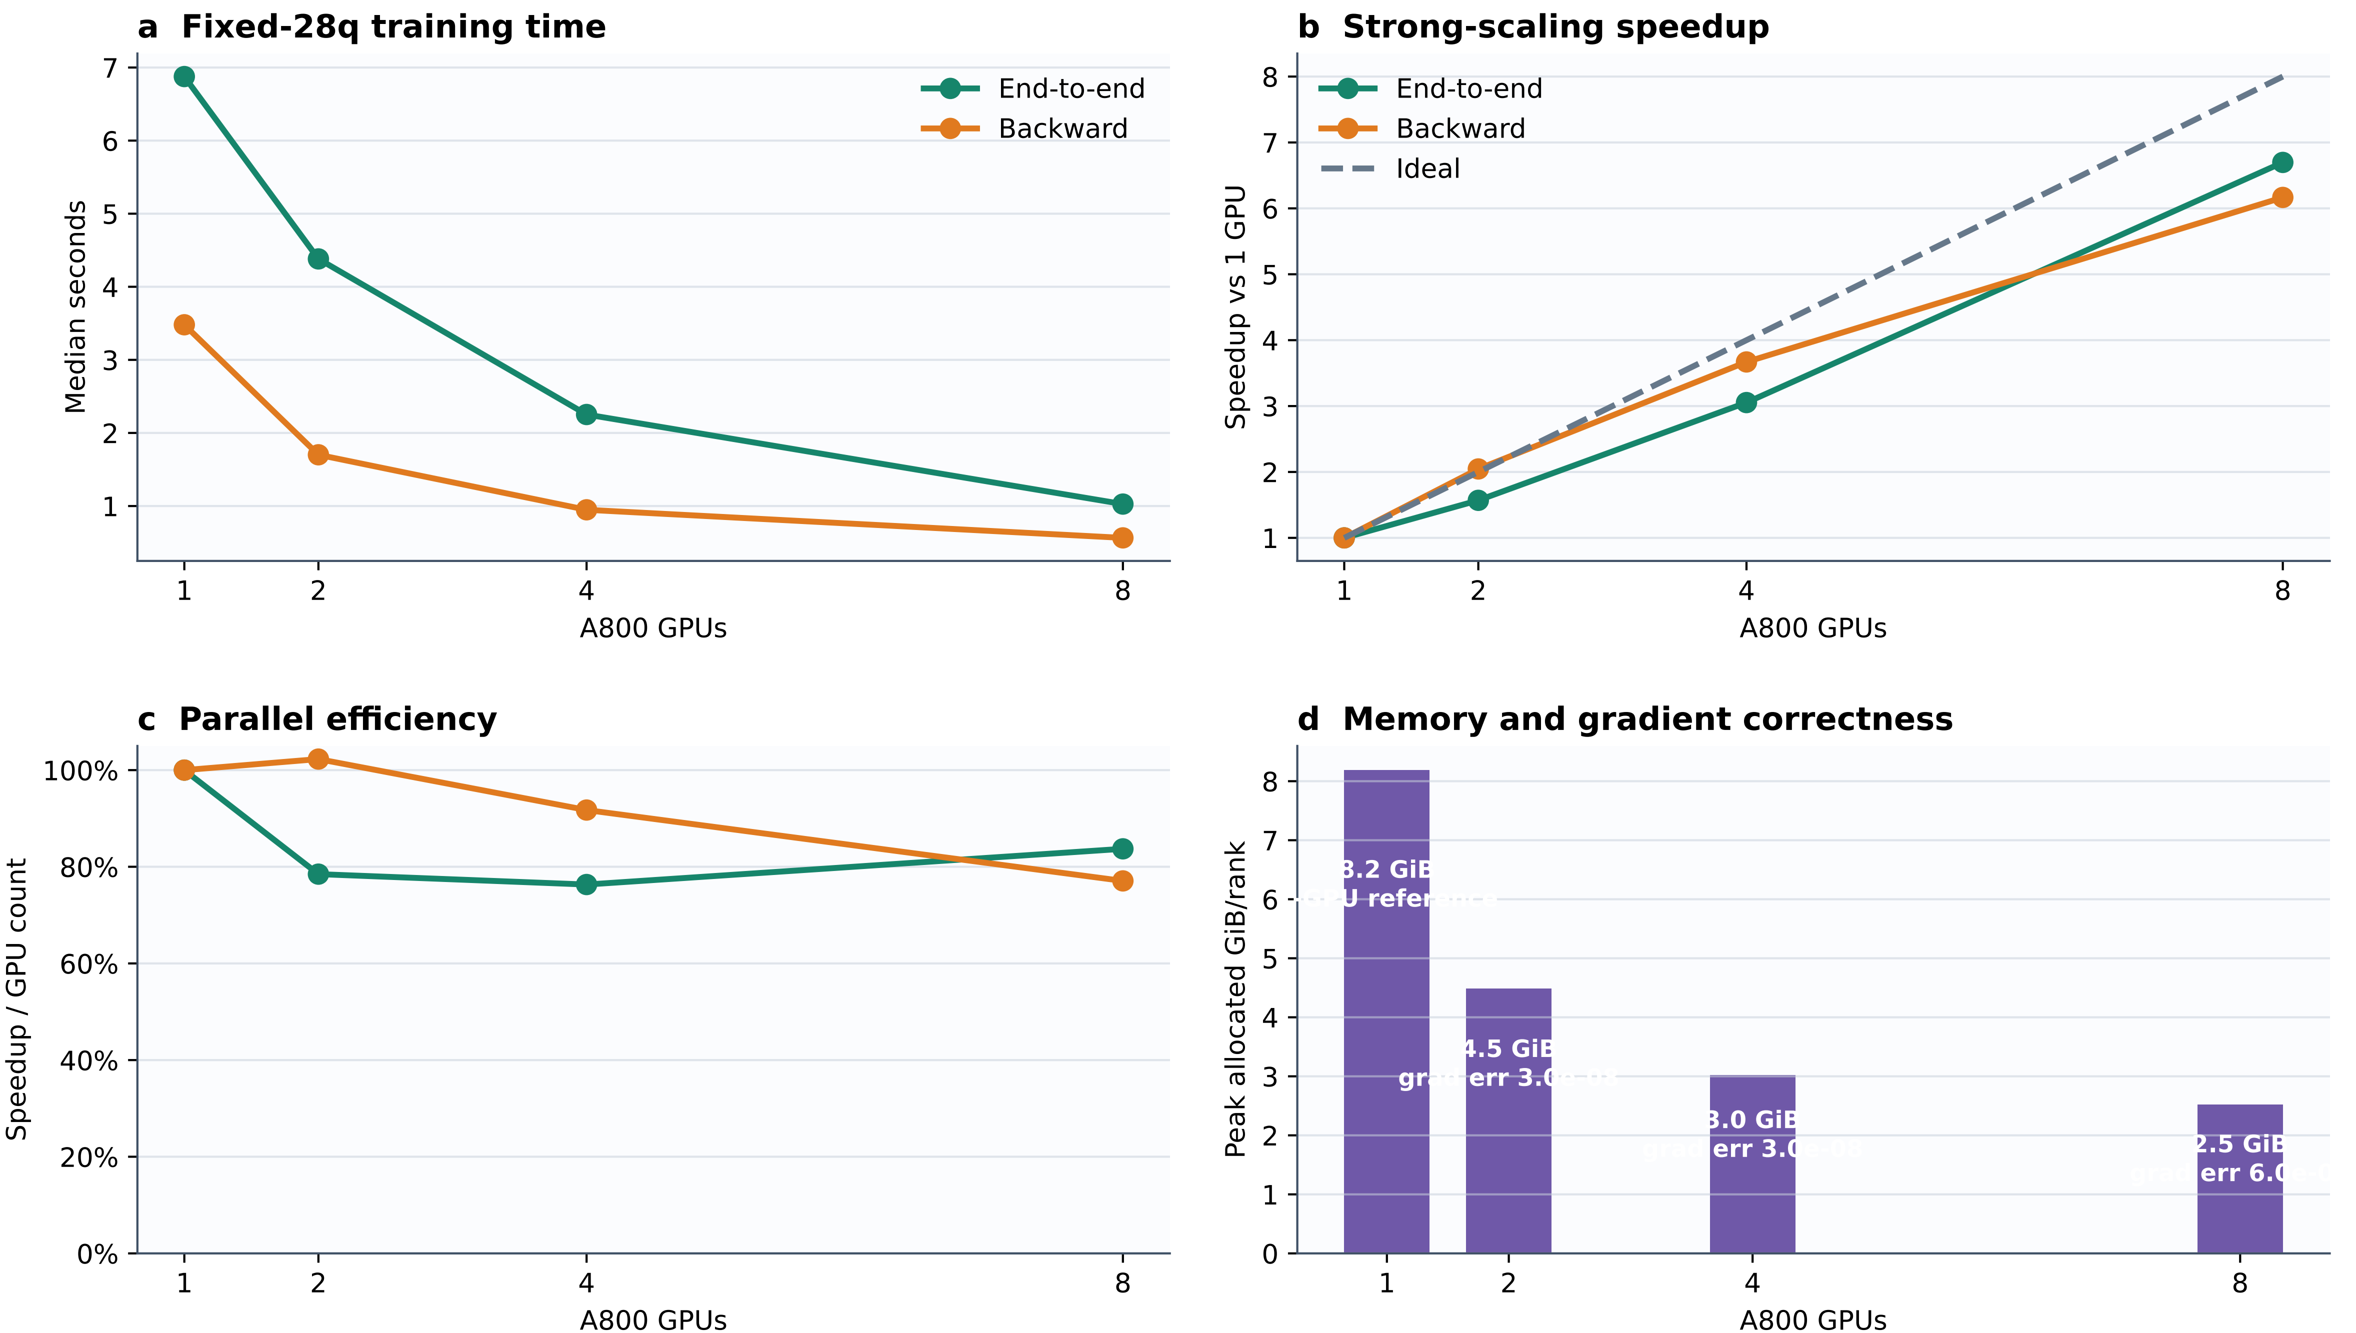

complex64 · 28 qubits · full-width linear HEA · RY + directed CX · sharded forward, adjoint backward and owner-sharded Adam

Where the speedup comes from · compute, communication and topology

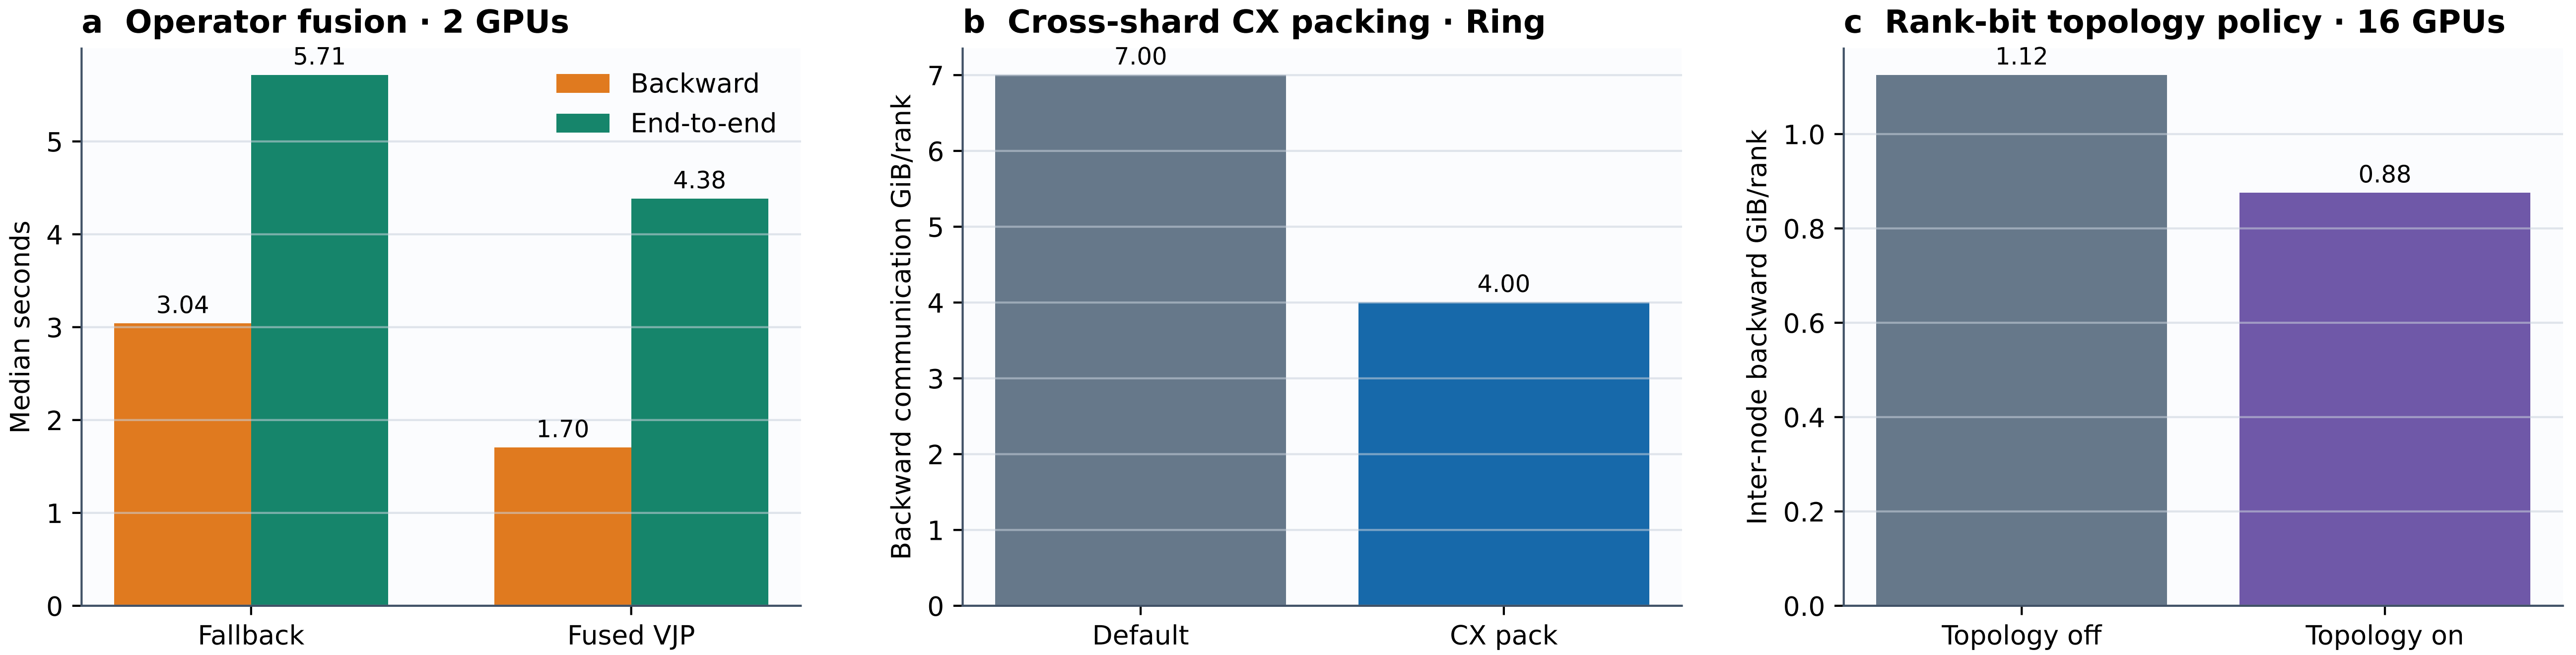

Optimization generality · Full-width Linear, Ring and Brickwork

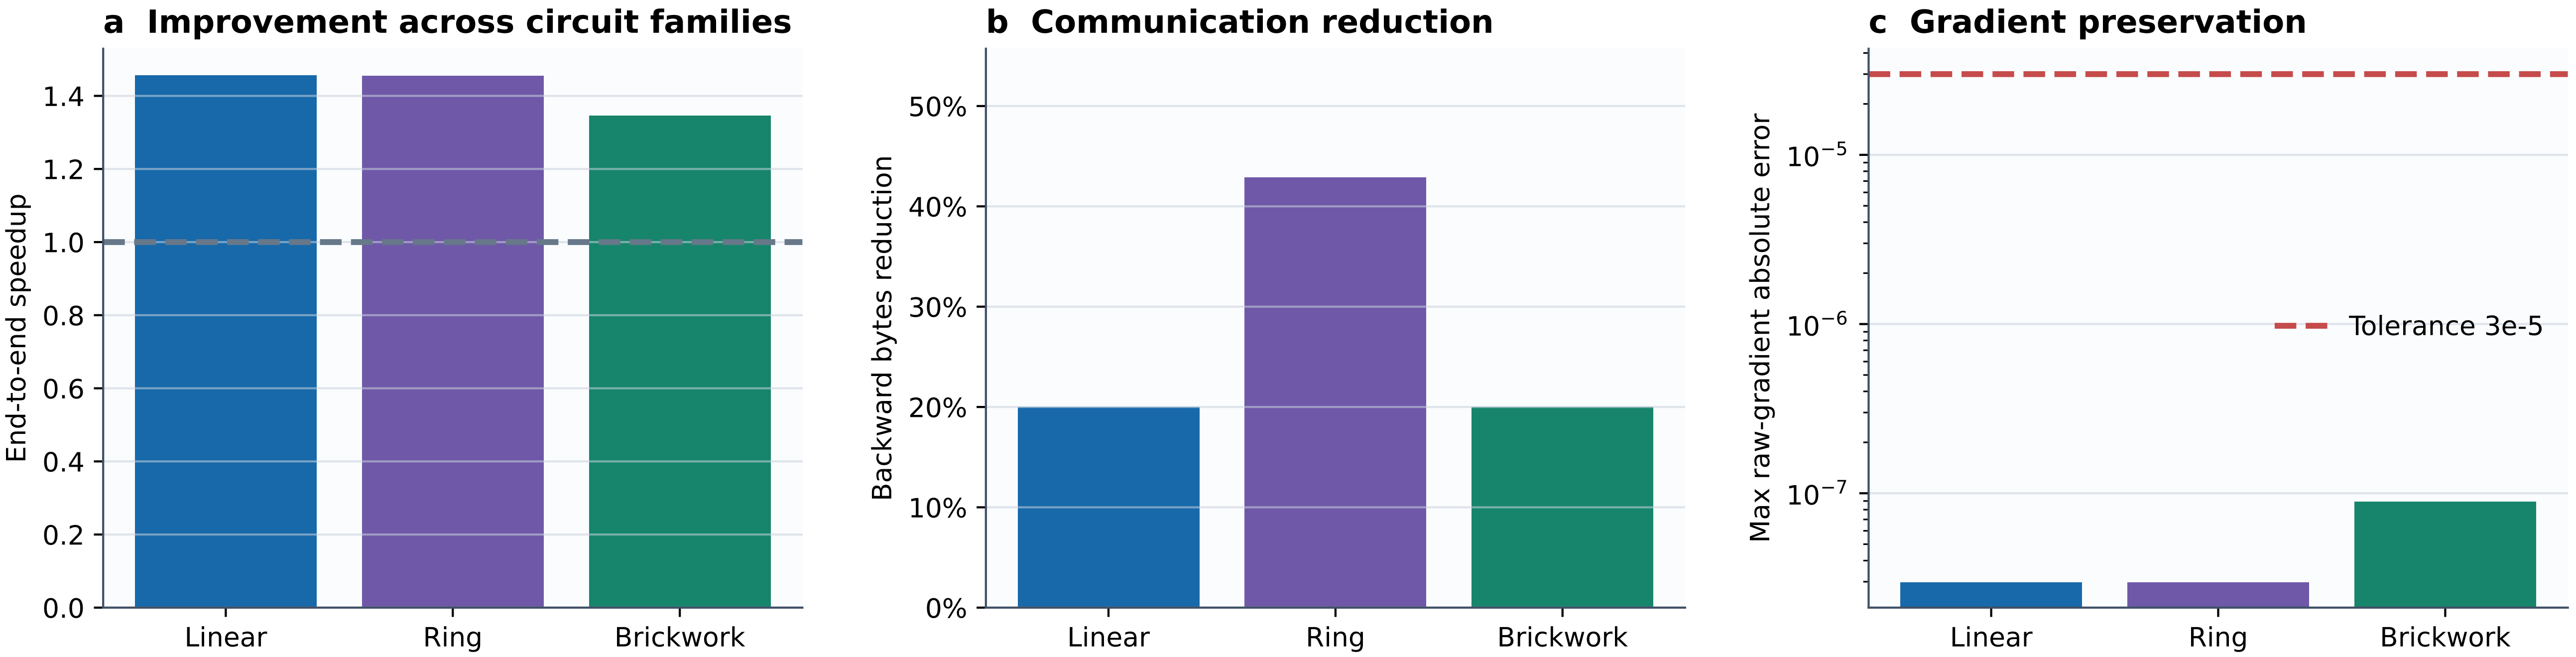

End-to-end result · faster training and a workload that only multi-GPU can run

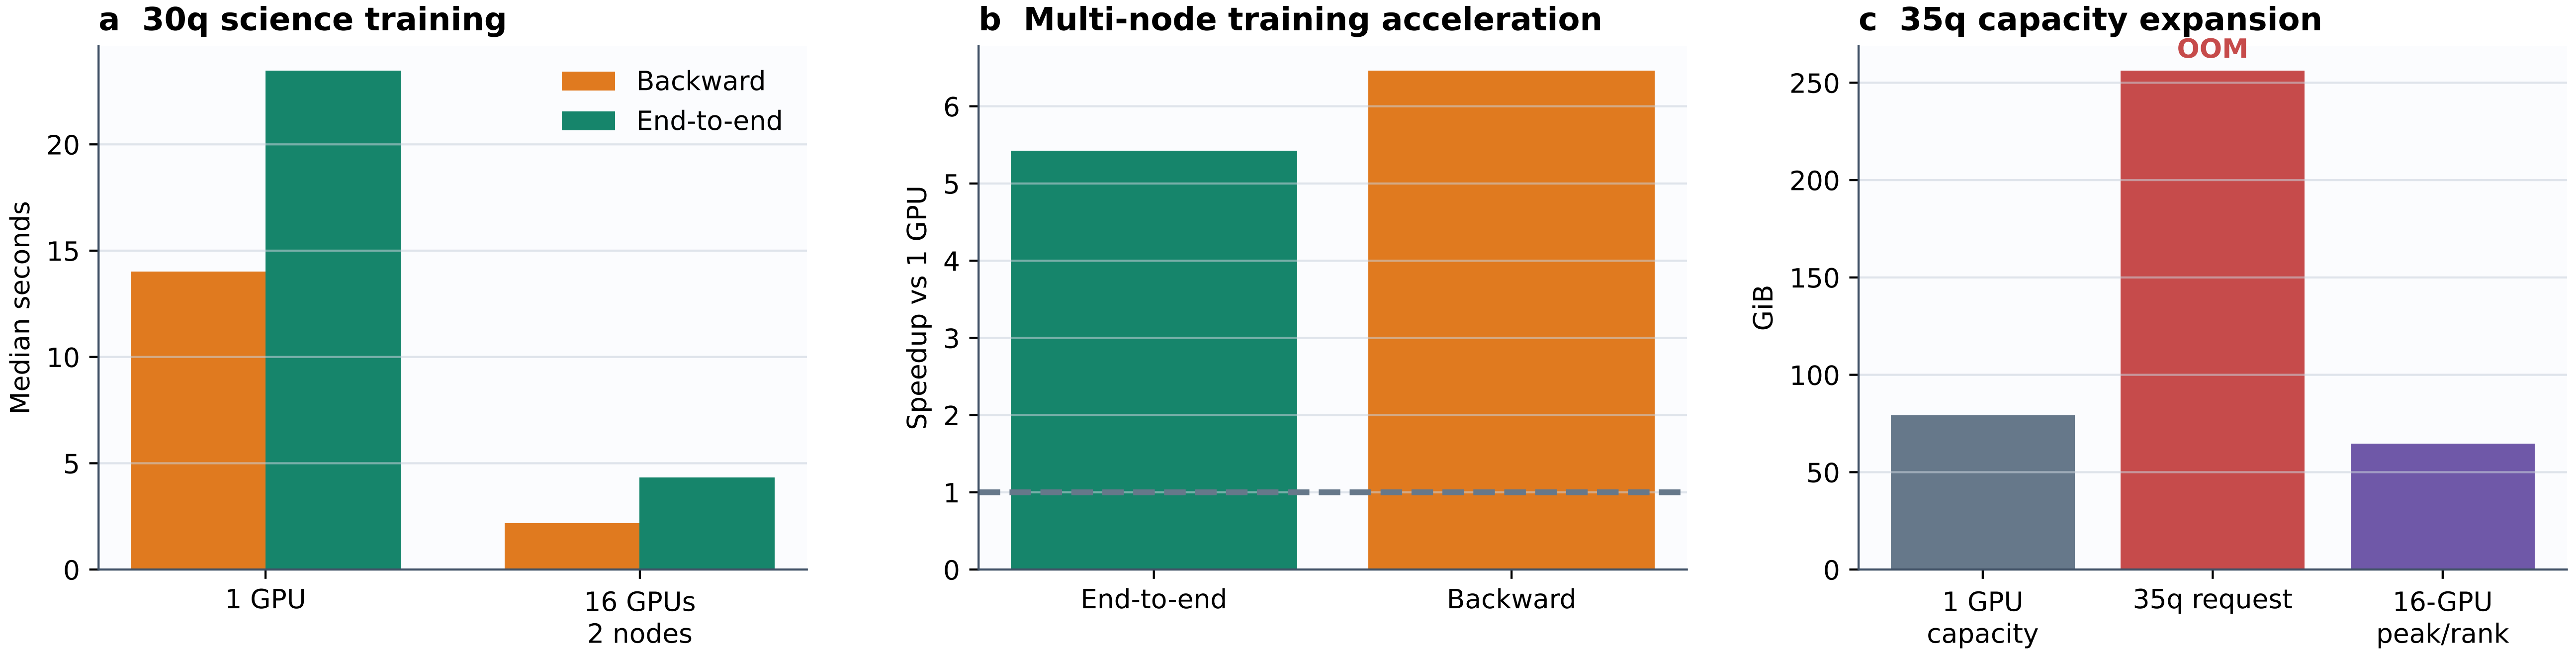

External black-box baseline · benchmark-only, no cuQuantum dependency

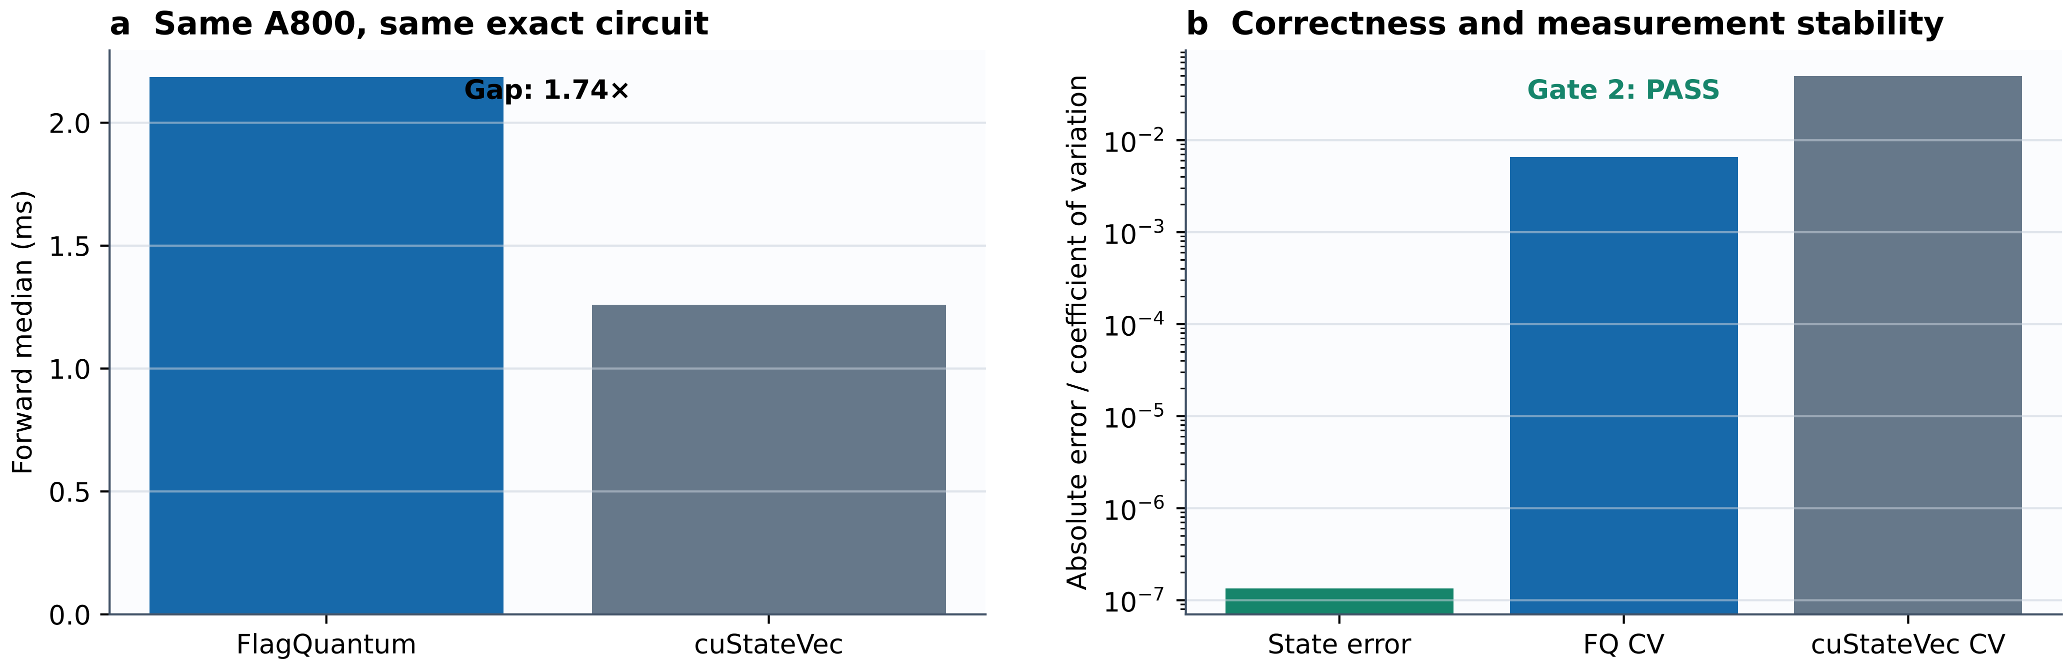

Frozen speed-sweep point · 24q, depth 8, 376 gates, 192 parameters

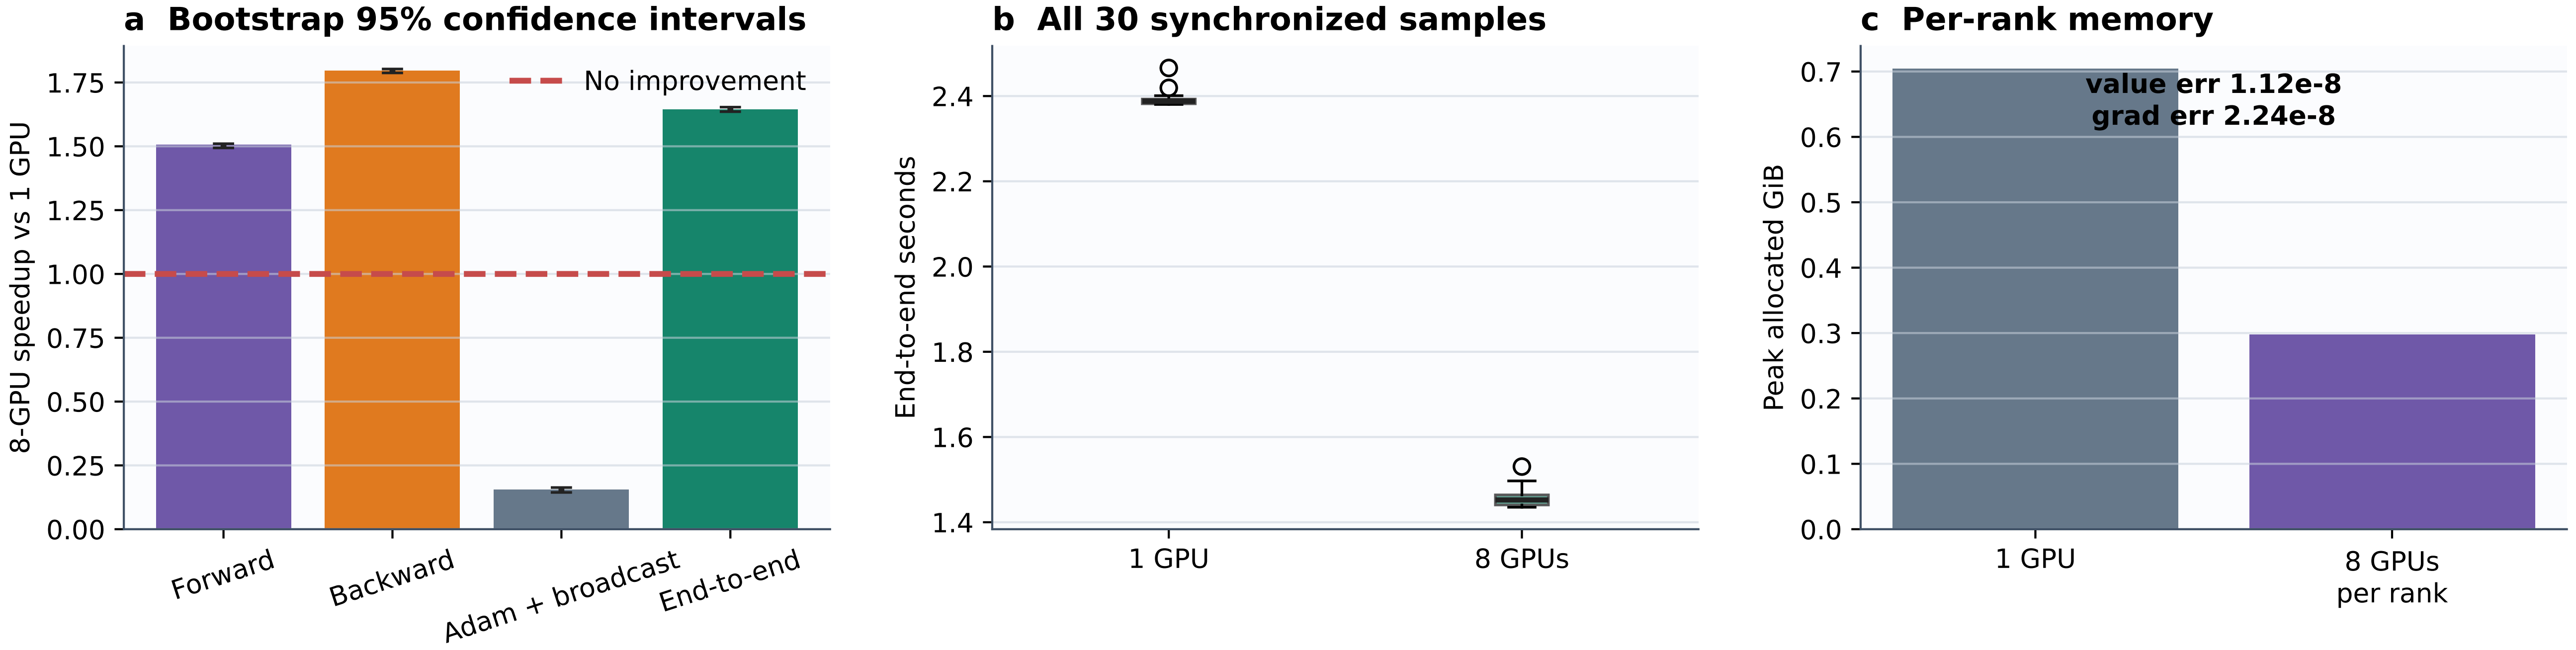

Owner-sharded optimizer synchronization · 192 broadcasts collapsed to one collective