



Matched differentiable full-width linear HEA

31 qubits · 8 layers · 248 trainable RY parameters · 488 gates


Directed nearest-neighbor CNOT chain · (Z_{15}) · complex64




a 31q end-to-end runtime



b Strong scaling



c Where framework time is recorded



d Relative performance

